

```
In [ ]: #Sentiment Classification Using BERT:
```

```
In [1]: #Step 1: Import the necessary Libraries:
```

```
import os
import shutil
import tarfile
import tensorflow as tf
from transformers import BertTokenizer, TFBertForSequenceClassification
import pandas as pd
from bs4 import BeautifulSoup
import re
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.offline as pyo
import plotly.graph_objects as go
from wordcloud import WordCloud, STOPWORDS
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
```

WARNING:tensorflow:From C:\Users\Ahmad\anaconda3\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
In [8]: #Step 2: Load the dataset:
```

```
In [3]: # Get the current working directory
```

```
current_folder = os.getcwd()

dataset = tf.keras.utils.get_file(
    fname="aclImdb.tar.gz",
    origin="http://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz",
    cache_dir=current_folder,
    extract = True)
```

```
In [4]: #check the dataset folder:
```

```
dataset_path = os.path.dirname(dataset)
# Check the dataset
os.listdir(dataset_path)
```

```
Out[4]: ['aclImdb', 'aclImdb.tar.gz']
```

```
In [7]: #Check the 'aclImdb' directory:
```

```
# Dataset directory
dataset_dir = os.path.join(dataset_path, 'aclImdb')

# Check the Dataset directory
os.listdir(dataset_dir)
```

```
Out[7]: ['imdb.vocab', 'imdbEr.txt', 'README', 'test', 'train']
```

```
In [8]: #Check the 'Train' dataset folder:
```

```
train_dir = os.path.join(dataset_dir, 'train')
os.listdir(train_dir)
```

```
Out[8]: ['labeledBow.feats',
        'neg',
        'pos',
        'unsup',
        'unsupBow.feats',
        'urls_neg.txt',
        'urls_pos.txt',
        'urls_unsup.txt']
```

```
In [9]: #Read the files of the 'Train' directory files:
```

```
for file in os.listdir(train_dir):
    file_path = os.path.join(train_dir, file)
    # Check if it's a file (not a directory)
    if os.path.isfile(file_path):
        with open(file_path, 'r', encoding='utf-8') as f:
            first_value = f.readline().strip()
            print(f"{file}: {first_value}")
    else:
        print(f"{file}: {file_path}")
```

```
labeledBow.feats: 9 0:9 1:1 2:4 3:4 4:6 5:4 6:2 7:2 8:4 10:4 12:2 26:1 27:1 28:1 29:2
32:1 41:1 45:1 47:1 50:1 54:2 57:1 59:1 63:2 64:1 66:1 68:2 70:1 72:1 78:1 100:1 106:
1 116:1 122:1 125:1 136:1 140:1 142:1 150:1 167:1 183:1 201:1 207:1 208:1 213:1 217:1
230:1 255:1 321:5 343:1 357:1 370:1 390:2 468:1 514:1 571:1 619:1 671:1 766:1 877:1 1
057:1 1179:1 1192:1 1402:2 1416:1 1477:2 1940:1 1941:1 2096:1 2243:1 2285:1 2379:1 29
34:1 2938:1 3520:1 3647:1 4938:1 5138:4 5715:1 5726:1 5731:1 5812:1 8319:1 8567:1 104
80:1 14239:1 20604:1 22409:4 24551:1 47304:1
neg: C:\Users\Ahmad\ML Lab\ML project 2.0\datasets\aclImdb\train\neg
pos: C:\Users\Ahmad\ML Lab\ML project 2.0\datasets\aclImdb\train\pos
unsup: C:\Users\Ahmad\ML Lab\ML project 2.0\datasets\aclImdb\train\unsup
unsupBow.feats: 0 0:8 1:6 3:5 4:2 5:1 7:1 8:5 9:2 10:1 11:2 13:3 16:1 17:1 18:1 19:1 2
2:3 24:1 26:3 28:1 30:1 31:1 35:2 36:1 39:2 40:1 41:2 46:2 47:1 48:1 52:1 63:1 67:1 6
8:1 74:1 81:1 83:1 87:1 104:1 105:1 112:1 117:1 131:1 151:1 155:1 170:1 198:1 225:1 2
26:1 288:2 291:1 320:1 331:1 342:1 364:1 374:1 384:2 385:1 407:1 437:1 441:1 465:1 46
8:1 470:1 519:1 595:1 615:1 650:1 692:1 851:1 937:1 940:1 1100:1 1264:1 1297:1 1317:1
1514:1 1728:1 1793:1 1948:1 2088:1 2257:1 2358:1 2584:2 2645:1 2735:1 3050:1 4297:1 5
385:1 5858:1 7382:1 7767:1 7773:1 9306:1 10413:1 11881:1 15907:1 18613:1 18877:1 2547
9:1
urls_neg.txt: http://www.imdb.com/title/tt0064354/usercomments
urls_pos.txt: http://www.imdb.com/title/tt0453418/usercomments
urls_unsup.txt: http://www.imdb.com/title/tt0018515/usercomments
```

```
In [10]: #Load the Movies reviews and convert them into the pandas' data frame with their respo
#Here 0 means Negative and 1 means Positive
```

```
def load_dataset(directory):
    data = {"sentence": [], "sentiment": []}
    for file_name in os.listdir(directory):
        print(file_name)
        if file_name == 'pos':
            positive_dir = os.path.join(directory, file_name)
            for text_file in os.listdir(positive_dir):
                text = os.path.join(positive_dir, text_file)
                with open(text, "r", encoding="utf-8") as f:
                    data["sentence"].append(f.read())
                    data["sentiment"].append(1)
            elif file_name == 'neg':
```

```

negative_dir = os.path.join(directory, file_name)
for text_file in os.listdir(negative_dir):
    text = os.path.join(negative_dir, text_file)
    with open(text, "r", encoding="utf-8") as f:
        data["sentence"].append(f.read())
        data["sentiment"].append(0)

return pd.DataFrame.from_dict(data)

```

In [11]: *#Load the training datasets*

```

# Load the dataset from the train_dir
train_df = load_dataset(train_dir)
print(train_df.head())

```

```

labeledBow.feats
neg
pos
unsup
unsupBow.feats
urls_neg.txt
urls_pos.txt
urls_unsup.txt

              sentence  sentiment
0  Story of a man who has unnatural feelings for ...      0
1  Airport '77 starts as a brand new luxury 747 p...      0
2  This film lacked something I couldn't put my f...      0
3  Sorry everyone,,, I know this is supposed to b...      0
4  When I was little my parents took me along to ...      0

```

In [12]: *#Load the test dataset respectively*

```

test_dir = os.path.join(dataset_dir, 'test')

# Load the dataset from the train_dir
test_df = load_dataset(test_dir)
print(test_df.head())

```

```

labeledBow.feats
neg
pos
urls_neg.txt
urls_pos.txt

              sentence  sentiment
0  Once again Mr. Costner has dragged out a movie...      0
1  This is an example of why the majority of acti...      0
2  First of all I hate those moronic rappers, who...      0
3  Not even the Beatles could write songs everyon...      0
4  Brass pictures (movies is not a fitting word f...      0

```

In [13]: *#Step 3: Preprocessing:*

```

In [14]: sentiment_counts = train_df['sentiment'].value_counts()

fig = px.bar(x= {0: 'Negative', 1: 'Positive'},
             y= sentiment_counts.values,
             color=sentiment_counts.index,
             color_discrete_sequence = px.colors.qualitative.Dark24,
             title='<b>Sentiments Counts')

```

```
fig.update_layout(title='Sentiments Counts',
                  xaxis_title='Sentiment',
                  yaxis_title='Counts',
                  template='plotly_dark')

# Show the bar chart
fig.show()
pyo.plot(fig, filename = 'Sentiments Counts.html', auto_open = True)
```

Out[14]: 'Sentiments Counts.html'

In [16]: *#Text Cleaning*

```
def text_cleaning(text):
    soup = BeautifulSoup(text, "html.parser")
    text = re.sub(r'\[[^\]]*\]', '', soup.get_text())
    pattern = r"^[a-zA-Z0-9\s,']"
    text = re.sub(pattern, '', text)
    return text
```

In [17]: *#Apply text_cleaning*

```
# Train dataset
train_df['Cleaned_sentence'] = train_df['sentence'].apply(text_cleaning).tolist()
# Test dataset
test_df['Cleaned_sentence'] = test_df['sentence'].apply(text_cleaning)
```

C:\Users\Ahmad\AppData\Local\Temp\ipykernel_18744\1023263618.py:4: MarkupResemblesLocatorWarning:

The input looks more like a filename than markup. You may want to open this file and pass the filehandle into Beautiful Soup.

C:\Users\Ahmad\AppData\Local\Temp\ipykernel_18744\1023263618.py:4: MarkupResemblesLocatorWarning:

The input looks more like a filename than markup. You may want to open this file and pass the filehandle into Beautiful Soup.


```
In [18]: #Plot reviews on WordClouds

# Function to generate word cloud
def generate_wordcloud(text, Title):
    all_text = " ".join(text)
    wordcloud = WordCloud(width=800,
                           height=400,
                           stopwords=set(STOPWORDS),
                           background_color='black').generate(all_text)

    plt.figure(figsize=(10, 5))
    plt.imshow(wordcloud, interpolation='bilinear')
    plt.axis("off")
    plt.title(Title)
    plt.show()
```

```
In [19]: #Positive Reviews


positive = train_df[train_df['sentiment']==1]['Cleaned_sentence'].tolist()
generate_wordcloud(positive, 'Positive Review')
```



```
In [20]: #Negative Reviews

negative = train_df[train_df['sentiment']==0]['Cleaned_sentence'].tolist()
generate_wordcloud(negative,'Negative Review')
```

Negative Review



```
In [21]: #Separate input text and target sentiment of both train and test
```

```
# Training data
#Reviews = "[CLS] " +train_df['Cleaned_sentence'] + "[SEP]"
Reviews = train_df['Cleaned_sentence']
Target = train_df['sentiment']

# Test data
#test_reviews = "[CLS] " +test_df['Cleaned_sentence'] + "[SEP]"
test_reviews = test_df['Cleaned_sentence']
test_targets = test_df['sentiment']
```

```
In [22]: #Split TEST data into test and validation
```

[illegible]

```
In [23]: #Step 4: Tokenization & Encoding
```

```
In [24]: #Load the pre-trained BERT tokenizer
```

```
#Tokenize and encode the data using the BERT tokenizer
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased', do_lower_case=True)
```

```
In [25]: #Apply the BERT tokenization in training, testing and validation dataset
```

```
max_len= 128
# Tokenize and encode the sentences
X_train_encoded = tokenizer.batch_encode_plus(Reviews.tolist(),
                                              padding=True,
                                              truncation=True,
                                              max_length = max_len,
                                              return_tensors='tf')

X_val_encoded = tokenizer.batch_encode_plus(x_val.tolist(),
```

```
padding=True,  
truncation=True,  
max_length = max_len,  
return_tensors='tf')  
  
X_test_encoded = tokenizer.batch_encode_plus(x_test.tolist(),  
padding=True,  
truncation=True,  
max_length = max_len,  
return_tensors='tf')
```

In [26]: *#Check the encoded dataset*

```
k = 0  
print('Training Comments -->', Reviews[k])  
print('\nInput Ids -->\n', X_train_encoded['input_ids'][k])  
print('\nDecoded Ids -->\n', tokenizer.decode(X_train_encoded['input_ids'][k]))  
print('\nAttention Mask -->\n', X_train_encoded['attention_mask'][k])  
print('\nLabels -->', Target[k])
```

Training Comments -->> Story of a man who has unnatural feelings for a pig Starts out with a opening scene that is a terrific example of absurd comedy A formal orchestra a udience is turned into an insane, violent mob by the crazy chantings of it's singers Unfortunately it stays absurd the WHOLE time with no general narrative eventually making it just too off putting Even those from the era should be turned off The cryptic dialogue would make Shakespeare seem easy to a third grader On a technical level it's better than you might think with some good cinematography by future great Vilmos Zsigmond Future stars Sally Kirkland and Frederic Forrest can be seen briefly

Input Ids -->>

```
tf.Tensor(
[[ 101  2466  1997  1037  2158  2040  2038 21242  5346  2005  1037 10369
   4627  2041  2007  1037  3098  3496  2008  2003  1037 27547  2742  1997
 18691  4038  1037  5337  4032  4378  2003  2357  2046  2019  9577  1010
   6355 11240  2011  1996  4689 22417  2015  1997  2009  1005  1055  8453
   6854  2009 12237 18691  1996  2878  2051  2007  2053  2236  7984  2776
   2437  2009  2074  2205  2125  5128  2130  2216  2013  1996  3690  2323
   2022  2357  2125  1996 26483  7982  2052  2191  8101  4025  3733  2000
   1037  2353  3694  2099  2006  1037  4087  2504  2009  1005  1055  2488
   2084  2017  2453  2228  2007  2070  2204 16434  2011  2925  2307  6819
 13728  2891  1062  5332 21693 15422  2925  3340  8836 11332  3122  1998
 15296 16319  2064  2022  2464  4780    102         0], shape=(128,), dtype=int32)
```

Decoded Ids -->>

[CLS] story of a man who has unnatural feelings for a pig starts out with a opening scene that is a terrific example of absurd comedy a formal orchestra audience is turned into an insane, violent mob by the crazy chantings of it's singers unfortunately it stays absurd the whole time with no general narrative eventually making it just too off putting even those from the era should be turned off the cryptic dialogue would make shakespeare seem easy to a third grader on a technical level it's better than you might think with some good cinematography by future great vilmos zsigmond future stars sally kirkland and frederic forrest can be seen briefly [SEP] [PAD]

Attention Mask -->>

```
tf.Tensor([1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0], shape=(128,), dtype=int32)
```

Labels -->> 0

```
In [27]: #Step 5: Build the classification model
```

```
In [28]: #Load the model
```

```
# Intialize the model
```

```
model = TFBertForSequenceClassification.from_pretrained('bert-base-uncased', num_label
```

```
model.safetensors: 0%|          | 0.00/440M [00:00<?, ?B/s]
```


C:\Users\Ahmad\anaconda3\Lib\site-packages\huggingface_hub\file_download.py:147: User Warning:

`huggingface_hub` cache-system uses symlinks by default to efficiently store duplicated files but your machine does not support them in C:\Users\Ahmad\.cache\huggingface\hub. Caching files will still work but in a degraded version that might require more space on your disk. This warning can be disabled by setting the `HF_HUB_DISABLE_SYMLINKS_WARNING` environment variable. For more details, see https://huggingface.co/docs/huggingface_hub/how-to-cache#limitations.

To support symlinks on Windows, you either need to activate Developer Mode or to run Python as an administrator. In order to see activate developer mode, see this article: <https://docs.microsoft.com/en-us/windows/apps/get-started/enable-your-device-for-development>

WARNING:tensorflow:From C:\Users\Ahmad\anaconda3\Lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

All PyTorch model weights were used when initializing TFBertForSequenceClassification.

Some weights or buffers of the TF 2.0 model TFBertForSequenceClassification were not initialized from the PyTorch model and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In [29]: *#Compile the model*

```
# Compile the model with an appropriate optimizer, loss function, and metrics
optimizer = tf.keras.optimizers.Adam(learning_rate=2e-5)
loss = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
metric = tf.keras.metrics.SparseCategoricalAccuracy('accuracy')
model.compile(optimizer=optimizer, loss=loss, metrics=[metric])
```

In [30]: *# Step 5: Train the model*

```
history = model.fit(
    [X_train_encoded['input_ids'], X_train_encoded['token_type_ids'], X_train_encoded['attention_mask']],
    Target,
    validation_data=(
        [X_val_encoded['input_ids'], X_val_encoded['token_type_ids'], X_val_encoded['attention_mask']],
        Target_val,
    ),
    batch_size=32,
    epochs=3
)
```

Epoch 1/3

WARNING:tensorflow:From C:\Users\Ahmad\anaconda3\Lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

782/782 [=====] - 9930s 13s/step - loss: 0.3477 - accuracy: 0.8410 - val_loss: 0.3082 - val_accuracy: 0.8734

Epoch 2/3

782/782 [=====] - 8050s 10s/step - loss: 0.2033 - accuracy: 0.9200 - val_loss: 0.2930 - val_accuracy: 0.8843

Epoch 3/3

782/782 [=====] - 8748s 11s/step - loss: 0.1078 - accuracy: 0.9606 - val_loss: 0.3566 - val_accuracy: 0.8864

In [31]: *#Step 6: Evaluate the model:*

```
In [32]: #Evaluate the model on the test data
test_loss, test_accuracy = model.evaluate(
    [X_test_encoded['input_ids'], X_test_encoded['token_type_ids'], X_test_encoded['at
    y_test
    )
print(f'Test loss: {test_loss}, Test accuracy: {test_accuracy}')
```

391/391 [=====] - 1083s 3s/step - loss: 0.3610 - accuracy: 0.8860
 Test loss: 0.360990047454834, Test accuracy: 0.8859999775886536

In [33]: *#Save the model and tokenizer to the local folder*

```
path = 'path-to-save'
# Save tokenizer
tokenizer.save_pretrained(path + '/Tokenizer')

# Save model
model.save_pretrained(path + '/Model')
```

In [34]: *#Load the model and tokenizer from the local folder*

```
# Load tokenizer
bert_tokenizer = BertTokenizer.from_pretrained(path + '/Tokenizer')

# Load model
bert_model = TFBertForSequenceClassification.from_pretrained(path + '/Model')
```

Some layers from the model checkpoint at path-to-save/Model were not used when initializing TFBertForSequenceClassification: ['dropout_37']

- This IS expected if you are initializing TFBertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing TFBertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

All the layers of TFBertForSequenceClassification were initialized from the model checkpoint at path-to-save/Model.

If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertForSequenceClassification for predictions without further training.

In [35]: *#Predict the sentiment of the test dataset*

```
pred = bert_model.predict(
    [X_test_encoded['input_ids'], X_test_encoded['token_type_ids'], X_test_encoded['at

# pred is of type TFSequenceClassifierOutput
logits = pred.logits

# Use argmax along the appropriate axis to get the predicted labels
pred_labels = tf.argmax(logits, axis=1)

# Convert the predicted labels to a NumPy array
pred_labels = pred_labels.numpy()

label = {
```

```

1: 'positive',
0: 'Negative'
}

# Map the predicted labels to their corresponding strings using the label dictionary
pred_labels = [label[i] for i in pred_labels]
Actual = [label[i] for i in y_test]

print('Predicted Label :', pred_labels[:10])
print('Actual Label :', Actual[:10])

391/391 [=====] - 1129s 3s/step
Predicted Label : ['Negative', 'positive', 'positive', 'Negative', 'Negative', 'Negative', 'Negative', 'positive', 'Negative', 'positive']
Actual Label : ['Negative', 'positive', 'positive', 'Negative', 'Negative', 'Negative', 'Negative', 'positive', 'Negative', 'Negative']

```

```

In [36]: #Classification Report

print("Classification Report: \n", classification_report(Actual, pred_labels))

```

```

Classification Report:

```

	precision	recall	f1-score	support
Negative	0.89	0.88	0.88	6250
positive	0.88	0.89	0.89	6250
accuracy			0.89	12500
macro avg	0.89	0.89	0.89	12500
weighted avg	0.89	0.89	0.89	12500

```

In [37]: #Step 7: Prediction with user inputs

```

```

In [38]: def Get_sentiment(Review, Tokenizer=bert_tokenizer, Model=bert_model):
# Convert Review to a list if it's not already a list
if not isinstance(Review, list):
    Review = [Review]

Input_ids, Token_type_ids, Attention_mask = Tokenizer.batch_encode_plus(Review,
                                                                           padding=Tr
                                                                           truncation
                                                                           max_length
                                                                           return_ter

prediction = Model.predict([Input_ids, Token_type_ids, Attention_mask])

# Use argmax along the appropriate axis to get the predicted labels
pred_labels = tf.argmax(prediction.logits, axis=1)

# Convert the TensorFlow tensor to a NumPy array and then to a list to get the pre
pred_labels = [label[i] for i in pred_labels.numpy().tolist()]
return pred_labels

```

```

In [39]: #Let's predict with our own review

```

```

Review = '''Bahubali is a blockbuster Indian movie that was released in 2015.
It is the first part of a two-part epic saga that tells the story of a legendary hero
The movie has received rave reviews from critics and audiences alike for its stunning
spectacular action scenes, and captivating storyline.'''
Get_sentiment(Review)

```

```
1/1 [=====] - 4s 4s/step  
Out[39]: ['positive']
```

```
In [ ]:
```